

19 June 2026

ModernMolBERT: A Compact SELFIES Encoder for Molecular Representation Learning

Jakob S. Madsen¹, Sara Angelucci¹, Christopher L. Porterfield¹, Xiaoting Liu¹, Alexander S. Hauser¹

1. Department of Drug Design and Pharmacology, University of Copenhagen

Abstract

We introduce ModernMolBERT, an encoder-only transformer model for small-molecule representation learning. ModernMolBERT adapts the ModernBERT architecture to molecular language modelling and pairs it with a SELFIES-aware Atom Pair Encoding (APE) tokeniser. Both released variants (small and base) are pre-trained with masked language modelling on approximately 2.4M unique SELFIES strings from ChEMBL 36 and evaluated as frozen molecular embedders on 25 binary classification benchmarks. ModernMolBERT-base attains a mean ROC-AUC of 77.9 (reported on a 0 to 100 scale), competitive with a strong fingerprint baseline and improving on comparable pre-trained string encoders by roughly five points. ModernMolBERT provides a compact, reproducible, open-source foundation for SELFIES-based molecular representation learning; we release the model weights, the APE tokeniser, and training code.

Keywords

molecular representation learning, molecular embeddings, chemical language models, SELFIES, masked language modelling, frozen embeddings, drug-discovery bench-marks, molecular encoder, SMILES

ModernMolBERT: A Compact SELFIES Encoder for Molecular Representation Learning

Jakob S. Madsen^{1,*}, Sara Angelucci¹, Christopher L. Porterfield¹, Xiaoting Liu¹, and Alexander S. Hauser¹

¹Center for Pharmaceutical Data Science Education, Department of Drug Design and Pharmacology, University of Copenhagen, Denmark

*Correspondence: jsmadsen@sund.ku.dk

Abstract

We introduce MODERNMOLBERT, an encoder-only transformer model for small-molecule representation learning. MODERNMOLBERT adapts the MODERNBERT architecture to molecular language modelling and pairs it with a SELFIES-aware *Atom Pair Encoding* (APE) tokeniser. Both released variants (*small* and *base*) are pre-trained with masked language modelling on approximately 2.4M unique SELFIES strings from ChEMBL 36 and evaluated as frozen molecular embedders on 25 binary classification benchmarks. MODERNMOLBERT-base attains a mean ROC-AUC of 77.9 (reported on a 0 to 100 scale), competitive with a strong fingerprint baseline and improving on comparable pre-trained string encoders by roughly five points. MODERNMOLBERT provides a compact, reproducible, open-source foundation for SELFIES-based molecular representation learning; we release the model weights, the APE tokeniser, and training code.

Keywords: molecular representation learning; molecular embeddings; chemical language models; SELFIES; masked language modelling; frozen embeddings; drug-discovery benchmarks; MODERNBERT.

1 Introduction

Molecular embeddings are the interface between molecular structure and modern machine learning methods. Embeddings are the vectors passed to property predictors, virtual-screening pipelines, similarity searches, clustering analyses, analogue-retrieval workflows, and dataset-triage tools. Their quality influences which chemical differences a model can see, which structure-activity patterns can be learned from sparse assay labels, and how well a representation transfers from known to unexplored chemical space.

Molecular representation learning is therefore not just a preprocessing choice but plays a key role in downstream performance. A good molecular embedder must preserve chemically meaningful structure, it must avoid artefacts introduced by molecular notation, and it should stay efficient enough to apply across large-scale compound libraries. The three design choices behind MODERNMOLBERT aim to address these needs. We use chemically valid SELFIES inputs, a SELFIES-aware tokeniser, and an efficient MODERNBERT encoder. This introduction explains and motivates each of these.

A long-standing way to represent a molecule is to map it to a fixed-length vector of hand-crafted descriptors, such as physicochemical properties and substructure-based fingerprints. The most prominent examples are circular fingerprints such as the Extended Connectivity Fingerprint (ECFP) [1], which hash the presence of circular atomic neighbourhoods at increasing radii into a fixed-length bit vector. ECFP4 in particular has been a workhorse of quantitative structure–activity relationship (QSAR) modelling. These fingerprints are compact and interpretable. However, their fixed dimensionality and bespoke design mean they cannot adapt to new structural motifs, lose information through hashing collisions, and cannot learn richer representations of structure–activity correlations from data. An early bridge between fixed fingerprints and fully learned representations was Mol2Vec [2], which applied the Word2Vec natural language algorithm to sequences of Morgan substructure identifiers, learning distributed embeddings for molecular fragments without any labelled data. Learnable alternatives that operate directly on molecular graphs, treating atoms as nodes and bonds as edges, have also been proposed [3, 4]. However, these require specialised aggregation schemes and sit outside the mainstream transformer pretraining paradigm.

The BERT [5] family of bidirectional encoder models showed that masked language model (MLM) pretraining on large unlabelled corpora produces representations that transfer to many downstream tasks with minimal fine-tuning. Here, molecules are represented as SMILES (Simplified Molecular-Input Line-Entry System) strings [6], a linear notation that encodes atomic connectivity as a depth-first traversal of the molecular graph. This sequential form is a natural substrate for the same sequence models originally developed for words and sentences.

SMILES molecular language models followed in rapid succession: ChemBERTa [7], MolBERT [8], ChemBERTa-2 [9], and MolFormer [10]. These models all established that MLM pretraining on SMILES yielded molecular representations competitive with fingerprint baselines on property-prediction benchmarks without task-specific feature engineering, and that this approach scales to billion-molecule corpora. Table 1 summarises their key design choices alongside MODERNMOLBERT; Section 2 discusses them in detail.

Model	Repr.	Tokeniser	Arch.	Params	Corpus	Size (M)
ChemBERTa [7]	SMILES	BPE	RoBERTa	5–46 [†]	PubChem	77
ChemBERTa-2 [9]	SMILES	BPE/atom	RoBERTa	5–46 [†]	PubChem	77
MolBERT [8]	SMILES	atom	BERT	85	ChEMBL	2
MolFormer [10]	SMILES	atom	Linear-attn.	46.8	ZINC+PubChem	>1,100
SElFormer [11]	SELFIES	BPE	RoBERTa	87	ChEMBL	2
ModernMolBERT-small	SELFIES	APE	ModernBERT	34.1	ChEMBL 36	2.4
ModernMolBERT-base	SELFIES	APE	ModernBERT	114.3	ChEMBL 36	2.4

Table 1: Overview of representative molecular string encoder models. *Repr.* (Representation): molecular string representation used for pretraining. *Tokeniser*: tokenisation strategy (BPE = Byte Pair Encoding; atom = atom-level; APE = Atom Pair Encoding). *Arch.* (Architecture): encoder backbone. *Params* (Parameters): trainable parameters in millions. *Size* (of dataset): pretraining corpus size in millions of molecules. Embedding dimensions for these and additional models are listed in Table 8 (Section D). **Bold** rows indicate the models introduced here. [†]The ChemBERTa family spans several released checkpoints of approximately 5–46M parameters; the 77M label denotes pretraining molecules, not parameters.

As a notation for language modelling, however, SMILES has two weaknesses. First, most

perturbations of a SMILES string are chemically invalid [12], so a model spends capacity learning to recognise ill-formed inputs. In addition, a single molecule can be represented by many equivalent strings, which inflates the effective vocabulary and complicates self-supervised objectives. Second, subword tokenisers such as Byte Pair Encoding (BPE) [13], applied naively to molecular strings, can split multi-character atomic symbols (e.g. C1 into C + 1) [14], and destroy the chemically meaningful units a molecular tokeniser should preserve.

SELFIES (Self-Referencing Embedded Strings) [12] addresses the validity problem by design. Its context-free grammar guarantees that every syntactically valid SELFIES string decodes to a chemically valid molecule under any editing or sampling operation, a guarantee SMILES cannot provide (Figure 2A). It removes invalid-string noise from the pretraining signal while preserving the sequential structure that makes transformer encoders effective. Yet SELFIES has been adopted mainly for molecular *generation*, while encoder-only SELFIES representation learning remains comparatively less explored [14]. To our knowledge the combination of SELFIES, a SELFIES-adapted atom-preserving tokeniser, and a MODERN-BERT-style encoder has not been systematically evaluated as a molecular embedder.

Tokenisation, the mapping from a raw string to discrete vocabulary units, is a central modelling choice. Atom-level tokenisation is chemically principled but yields long sequences and misses recurring substructures, whereas BPE operates on raw character frequencies and produces entries that straddle atom boundaries. Leon et al. [14] proposed *Atom Pair Encoding* (APE), which adapts the BPE merge procedure to respect molecular structure by restricting merges to adjacent complete molecular symbols. We adapt APE to SELFIES by merging only adjacent complete SELFIES primitives, so every vocabulary entry is a valid concatenation of complete SELFIES symbols rather than a character-level artefact (Figure 2C). APE thus occupies a useful middle ground: compact like subword tokenisation, chemically principled like atom-level tokenisation.

Many molecular string encoders such as ChemBERTa, ChemBERTa-2, MolBERT, and SELFormer are built on the BERT- or RoBERTa-era backbones [5, 15] with absolute positional embeddings and full quadratic attention. Several newer post-BERT advances have been combined in MODERNBERT [16]: Rotary Position Embeddings (RoPE), Flash Attention, alternating local/global attention, and GeGLU activations (Section 2.3). To our knowledge, these innovations together have not been applied to SELFIES-based molecular language modelling. RoPE in particular suits molecular strings, where absolute sequence position is partly an artefact of the chosen graph traversal rather than an intrinsic molecular property. Because RoPE conditions attention on the *relative* offset between two tokens rather than their absolute indices, it may reduce the encoder’s sensitivity to where a traversal happens to begin. This starting point is one part of SMILES/SELFIES positioning that carries no chemical meaning.

The masked language model (MLM) objective does not directly specify which tokens to mask. Standard BERT masks tokens independently at random [5], whereas SpanBERT [17] masks contiguous spans to encourage modelling of longer-range structure. Masking can also be made chemically informed: Peng et al. [18] mask complete functional groups identified by parsing, reasoning that functional groups carry disproportionate property-relevant information. Whether such structure-aware masking benefits a frozen SELFIES embedder is an open question we examine in an ablation (Section 5.2).

We introduce MODERNMOLBERT, a compact encoder-only molecular language model that brings together three design choices that have not previously been evaluated together: chemically valid SELFIES inputs, a SELFIES-adapted APE tokeniser, and a MODERNBERT

backbone. The models are trained from scratch with MLM on 2.4M curated ChEMBL 36 molecules and released in small and base variants. We evaluate them as frozen molecular embedders on the 25-dataset benchmark of Praski et al. [19] against a fingerprint baseline and representative pre-trained string encoders. Figure 1 shows the pretraining and frozen-evaluation pipeline, and Figure 2 illustrates the two molecular design choices that distinguish MODERNMOLBERT from prior SELFIES encoders.

2 Related Work

The introduction motivated the three design choices behind MODERNMOLBERT; this section places them in the context of prior work along the four axes most relevant to a frozen SELFIES embedder: molecular string representations, molecular language models, modern encoder architectures, and molecular tokenisation. It also identifies the specific models we later use as baselines.

2.1 Molecular String Representations

The chemical-validity guarantee of SELFIES holds under arbitrary mutation. As reported by Krenn et al. [12], a single random bit-flip leaves a SMILES string valid only 26.6% of the time and ten bit-flips reduce validity to 0.2%, whereas SELFIES remains 100% valid in both cases. SELFIES is also complete (every molecule can be encoded) and serves as a drop-in input to any model. SELFIES thus offers a chemically attractive input space, yet the field still lacks a compact, modern encoder that tests whether SELFIES strings can serve as strong off-the-shelf molecular embeddings under a broad frozen benchmark.

2.2 Molecular Language Models

The fingerprint ECFP [1] remains the dominant deterministic baseline: a circular-substructure fingerprint with no trainable parameters that competes handily with learned representations on many property prediction tasks. Among learned string encoders, ChemBERTa-2 [9] revisited the RoBERTa-on-SMILES recipe, pretraining on PubChem SMILES and comparing masked language modelling with multi-task regression (MTR) over computed RDKit descriptors. MolFormer [10] scaled the masked-language-modelling approach to 1.1B molecules and introduced a linear-attention transformer, making it the largest-scale SMILES string model and a key large-scale reference point. SELFormer [11] is the closest prior work to MODERNMOLBERT: a RoBERTa encoder pre-trained with MLM on 2M drug-like SELFIES strings from ChEMBL. Several adjacent molecular foundation models use different pretraining objectives or architectures. Chemformer [20] uses a sequence-to-sequence transformer for computational chemistry tasks, whereas CDDD [21] learns continuous molecular descriptors by translating between equivalent molecular strings. Recent encoder-only work such as ChemBERTa-3 [22] and MolEncoder [23] continues to refine training recipes for chemical BERT-style models. Recent contrastive and multimodal approaches such as CLAMP [24] and COATI [25] extend molecular embeddings beyond pure string reconstruction by incorporating assay text, 3D information, or other paired supervision. These models broaden the representation-learning landscape but answer a different question from the one addressed here: how far a compact, encoder-only SELFIES MLM can be pushed as a frozen molecular embedder. A large and growing literature applies graph neural networks and graph transformers to molecular representation learning; representative examples include GROVER [4],




Figure 1: **ModernMolBERT pretraining and frozen-evaluation pipeline.** (A) Approximately 2.4M unique ChEMBL 36 small molecules are curated, canonicalised, and converted to SELFIES strings. (B) The SELFIES strings are pre-tokenised with the SELFIES-adapted APE tokeniser (Figure 2C). (C) A MODERNBERT encoder is trained from scratch on this corpus with the masked language modelling (MLM) objective, yielding MODERNMOLBERT. (D) For evaluation the encoder is *frozen*: mean-pooled final-layer token embeddings over non-special SELFIES tokens are extracted without any gradient updates and passed to lightweight downstream classifiers (ridge, random forest, k -nearest neighbours) on 25 binary molecular property-prediction benchmarks (7 biophysical endpoints from MoleculeNet and 18 toxicity, ADME and HTS endpoints from the Therapeutics Data Commons), where MODERNMOLBERT is compared against fixed fingerprint and pre-trained string-encoder baselines under an identical protocol.




Figure 2: **Method comparison of the two molecular design choices in ModernMol-BERT: input representation and tokenisation.** (A) *Input representation: SMILES versus SELFIES.* The same molecule is shown encoded as a SMILES string (top) and as a SELFIES string (bottom). SMILES tokens can violate chemical syntax under perturbation or naive tokenisation (mismatched brackets and ill-formed ring-closure indices are marked with red crosses), whereas every syntactically valid SELFIES *string* decodes to a chemically valid molecule [12]. (B) *Byte Pair Encoding (BPE).* BPE merges SELFIES symbols by raw character frequency without chemical awareness, splitting bracketed symbols and producing chemically meaningless vocabulary entries (dashed red markers). (C) *Atom Pair Encoding (APE).* APE restricts merges to adjacent complete SELFIES primitives, so every vocabulary entry is a valid concatenation of complete SELFIES symbols rather than a character-level artefact [14].

GIN [26], GraphMVP [27], and Uni-Mol [28]. These graph-based methods are outside the scope of MODERNMOLBERT, which operates on sequential SELFIES strings without explicit graph structure. Broader benchmarking studies have nevertheless shown that the relative value of learned molecular representations depends strongly on the dataset, split strategy, chemical-series structure, and amount of labelled data available [29, 30]. This motivates treating ECFP4 and other fixed representations as serious baselines rather than as outdated reference points. Collectively these models show the value of molecular pretraining, but leave open whether a smaller, validity-preserving encoder can supply a useful default representation without fine-tuning or specialised runtime dependencies.

2.3 Modern BERT Architectures

MODERNBERT [16] is an encoder-only transformer that consolidates post-BERT architectural advances into a production-quality model trained on 2 trillion tokens of text, code, and scientific literature with a native 8192-token context window. Its key design choices are: Rotary Position Embeddings (RoPE) [31] replacing absolute positional encodings; *alternating local and global attention*, in which most layers apply a sliding-window kernel and periodically interleaved layers apply full-sequence attention, executed with the IO-aware Flash Attention [32] kernel; GeGLU feed-forward activations; and unpadded sequence processing, which removes padding tokens to eliminate wasted compute. The two released variants for natural language use are MODERNBERT-base (149M parameters) and MODERNBERT-large (395M). They are trained with a 30% MLM masking rate and achieve state-of-the-art results among encoder-only models on natural-language and code benchmarks. MODERNMOLBERT tests whether these encoder improvements transfer to molecular SELFIES MLM without requiring a very large model or a non-standard downstream workflow.

2.4 Molecular Tokenisation

The atom-preserving APE tokeniser [14] (Section 3.1) was shown to outperform standard BPE in a controlled comparison of RoBERTa-based models pre-trained on 1M PubChem molecules and fine-tuned on MoleculeNet, with the largest gains on SELFIES representations (ROC-AUC improvements of +0.059 on BBBP and HIV and +0.017 on Tox21; Table 6 of Leon et al. 14), a result attributed to its merges aligning with SELFIES’s symbol-level structure. We therefore treat APE as a practical way to compress recurring SELFIES patterns while preserving a SELFIES-based input space, and ask whether this design helps empirically rather than assuming it is chemically optimal.

3 Methods

3.1 Atom Pair Encoding Tokeniser

We use a custom reimplementaion of the APE algorithm [14, 33] with several modifications for SELFIES pretraining at scale, detailed in the paragraphs below: molecular-aware pre-tokenisation that restricts merges to complete SELFIES primitives, frequency- and span-controlled vocabulary construction, deterministic greedy longest-match inference, and post-training injection of benchmark-observed primitives.

Molecular-aware pre-tokenisation. SELFIES strings are first split on a regex that isolates each complete bracket-enclosed primitive, so merges operate only over whole SELFIES primitives and every vocabulary entry is a valid concatenation of complete symbols.

Vocabulary construction. Merging is controlled by two criteria: a minimum merge frequency (`min_freq` = 3000), which halts merging once the most frequent adjacent pair co-occurs fewer than 3000 times, and a maximum merge span (`max_merge_pieces` = 2), which excludes candidates spanning more primitives to avoid long, uninformative concatenations. In practice `min_freq` is the binding early-stopping criterion. Full tokeniser statistics are reported in Table 5 (Section B).

Greedy longest-match tokenisation. At inference, each position is consumed by the longest matching vocabulary entry, looking up to `max_merge_pieces` positions ahead. The procedure is deterministic and $O(n \cdot \text{max_merge_pieces})$ per sequence, with an `<unk>` fallback only when no single primitive is in-vocabulary.

Post-training symbol injection. After merge learning but before model pretraining, any valid SELFIES primitive that appears ≥ 10 times in the benchmark molecules but is absent from the learned vocabulary is appended. This uses only primitive-frequency information; no benchmark labels or train/test split information are used in this step. It guarantees a zero unknown-token rate on evaluation molecules; all model weights were initialised and trained on the completed 631-token vocabulary from the outset. This step prevents out-of-vocabulary failures rather than providing task information: it mirrors the fixed, closed vocabulary that any deployed encoder must define before use. No labels or train/test split assignments enter the vocabulary, and the completed vocabulary is fixed before any model weight is trained, so the injected primitives carry no predictive signal about the downstream tasks.

3.2 Model Architecture

Both MODERNMOLBERT variants use the MODERNBERT encoder backbone (Section 2.3) [16] in its default configuration: RoPE, GeGLU feed-forward blocks, pre-normalisation, no linear-layer biases. Both use unpadded Flash Attention, with global attention every three layers and a 128-token local sliding window. The variants are trained from scratch on our SELFIES APE vocabulary; no weights are transferred from the official MODERNBERT checkpoints [34]. Per-variant hyper-parameters are listed in Table 3 (Section A). MODERNMOLBERT-small (34.1M parameters, 512-dimensional hidden state) sits in the compact regime below Chemformer [20] (45M), MoLFormer-XL [10] (46.8M), and Uni-Mol [28] (47.6M), and is notably smaller than the closest SELFIES comparator SELFormer (86.7M). MODERNMOLBERT-base (114.3M parameters, 768-dimensional) is substantially larger than the other models considered here.

3.3 Pretraining Data

Pretraining data are sourced from ChEMBL 36 [35] via the lukaskim/ChEMBL-36 Hugging Face dataset¹. Molecules are deduplicated by `standard_inchi_key`, parsed and lightly sanitised with RDKit [36] (ring perception and aromaticity assignment only, avoiding full sanitisation), and canonicalised to isomeric canonical SMILES using RDKit’s `MolToSmiles`

¹<https://huggingface.co/datasets/lukaskim/ChEMBL-36>

with `isomericSmiles=True`. The resulting canonical SMILES strings are then encoded to SELFIES [12] using the `selfies` Python package. Molecules that fail RDKit parsing, limited sanitisation, canonicalisation, or SELFIES encoding are excluded. We further filter by size and type, retaining only entries flagged as small molecules with $3 \leq \text{heavy atoms} \leq 100$ and molecular weight ≤ 1000 Da. Train and validation splits (99%/1%) are assigned deterministically by SHA-1 hashing of a stable per-molecule identifier (preferring the standard InChIKey, with fallbacks to ChEMBL ID, canonical SMILES, and SELFIES in that order) with seed 13, yielding 2,390,317 training and 24,228 validation molecules. The resulting SELFIES strings are subsequently pre-tokenised with the APE tokeniser [14] using greedy longest-match tokenisation. The full curation and pre-tokenisation pipeline is available at <https://github.com/HauserGroup/ModernMolBERT>. The final dataset is available on <https://huggingface.co/HauserGroup>.

3.4 Pretraining Procedure

Both MODERNMOLBERT variants are trained from scratch with the masked language modelling (MLM) objective and no auxiliary tasks. Training runs for 30,000 steps at an effective batch size of 256 and a maximum sequence length of 128 tokens. For each variant, the masking rate and peak learning rate are selected by a 3×3 grid search (Section 5.2); the remaining optimisation settings (AdamW, weight decay, gradient clipping, bf16, linear warmup, and cosine decay) are fixed and reported in Table 4 (Section A). Validation loss and masked-token accuracy converge smoothly within the fixed step budget (Figure 8).

Masking strategies. We implement three masking strategies as a selectable data-collator component. *Standard* masking draws an independent Bernoulli mask per eligible token. *Span* masking masks contiguous spans with lengths drawn from a Geometric distribution ($p = 0.4$, clamped to 6 tokens), following SpanBERT’s span selection but without its Span Boundary Objective [17]. *Heteroatom-biased span* masking up-weights span-start positions at APE tokens containing a heteroatom primitive (N, O, S, P, F, Cl, Br, I, Se, Si), identified by a bracket-aware SELFIES regex; spans extend beyond the start, so masking is enriched around rather than restricted to these tokens. It is a parser-free proxy for the functional-group masking of Peng et al. [18], using the APE bracket vocabulary in place of explicit SMARTS matching. All three apply the standard BERT corruption rule (80% [MASK], 10% random token, 10% unchanged). The released models use standard masking; Section 5.2 compares all three.

3.5 Evaluation Protocol

All benchmark tasks are evaluated using frozen MODERNMOLBERT embeddings, following the lightweight protocol of Praski et al. [19].² For each molecule, the encoder is run in inference mode and molecule-level representations are obtained by mean-pooling the final-layer token embeddings over non-special SELFIES tokens; no model parameters are updated. These embeddings are evaluated with three lightweight downstream models chosen to span complementary inductive biases: a linear ridge classifier, ensemble random forests, and instance-based k -nearest neighbours. This way, representation quality is not judged through a single decision boundary. We adopt these three heads, and their cross-validation search

²Using their public benchmark code at https://github.com/MLCIL/benchmarking_molecular_models.

grids, directly from the benchmark protocol of Praski et al. [19] to keep the comparison with their reported baselines exact; hyperparameters are selected by cross-validation on the training split, with performance reported on the held-out test split. All tasks are scored by ROC-AUC. For multi-task datasets, missing endpoint labels are treated as unmeasured: downstream classifiers are fit endpoint-wise on finite labels (in both cross-validation and test evaluation), and dataset-level ROC-AUC is averaged over endpoints with finite labels and both classes present in the test split. The best downstream model per dataset is used in the main comparison.

3.6 Design Rationale

Several configuration choices admit reasonable alternatives; we record the reasoning here rather than leave it implicit. We use mean pooling over non-special token embeddings, rather than a dedicated [CLS] token, because the evaluation is frozen: with no fine-tuning to adapt a special token to each task, a simple architecture-agnostic pool over all non-special tokens is a more robust molecule-level summary. The maximum sequence length of 128 tokens was set because the SELFIES+APE representation covers essentially the entire filtered ChEMBL 36 corpus, with fewer than 0.1% of molecules truncated at this length (Section B). The APE thresholds (`min_freq` = 3000, `max_merge_pieces` = 2) were chosen to give moderate compression while avoiding long, sparse merges; they are pragmatic engineering settings, not a claim that this tokenisation is chemically optimal. The 30,000-step budget was chosen as a reproducible medium-scale regime, with the validation curves (Figure 8) confirming that the selected runs were stable and had reached a performance plateau.

4 Experiments

4.1 Benchmark Datasets

We evaluate on 25 binary classification datasets drawn from two sources: 18 datasets from the Therapeutics Data Commons (TDC) [37], spanning toxicity (4 datasets), ADME (12 datasets), and high-throughput screening (2 datasets); and 7 datasets from MoleculeNet [38], covering a range of biophysical and physical chemistry endpoints. Dataset sizes range from 475 to 93 087 compounds, with task counts ranging from 1 to 27 (SIDER). The full list with per-dataset statistics is given in Table 4 of Praski et al. [19], whose benchmark selection we follow directly. All datasets are evaluated by ROC-AUC.

4.2 Baselines

We compare MODERNMOLBERT against both fixed and learned molecular representations. ECFP4 is the classical non-neural baseline [1, 29, 30]. The pre-trained string-model baselines each isolate one axis of comparison: ChemBERTa-2 is the closest-scale RoBERTa-style SMILES MLM; SELFFormer isolates the use of SELFIES with a BERT/RoBERTa-era encoder; and MoLFormer³ is a large-scale SMILES reference trained on substantially broader PubChem and ZINC data. Because MODERNMOLBERT is trained purely with MLM, we use the MLM-trained variant of each baseline so that the comparison is not confounded by the pretraining

³MoLFormer’s inference relies on the ONNX Runtime for its linear-attention kernel, which has limited compatibility with Python newer than ≈ 3.10 ; embedding extraction therefore required a pinned legacy environment.

objective: SELFormer and MoLFormer are MLM-pre-trained [10, 11], and for ChemBERTa-2 we use the MLM checkpoint (seyonec/ChemBERTa-77M-MLM) rather than the multi-task-regression variant [9]. Auxiliary-objective and multi-task-regression pretraining therefore falls outside this MLM-only comparison. No model, including MODERNMOLBERT, is fine-tuned end-to-end; all neural baselines are evaluated as frozen embedders under the same downstream protocol.

5 Results

5.1 ModernMolBERT Approaches Fingerprints and Improves on Closest-Scale String Encoders

The main comparison aggregates ROC-AUC across the 25 binary classification benchmark datasets described in Section 4.1. For each representation–dataset pair, we report the best cross-validated downstream estimator among ridge classification, random forests, and k -nearest neighbours, following the frozen mean-pooled embedding protocol in Section 3.5. For multi-endpoint datasets, endpoint ROC-AUCs are averaged before aggregation. This presentation separates the quality of the pre-trained representation from gains attributable to end-to-end fine-tuning, and therefore directly tests whether MODERNMOLBERT provides useful off-the-shelf molecular embeddings. Throughout, ROC-AUC values are reported $\times 100$, so one point corresponds to 0.01 ROC-AUC. Direct paired comparisons use only jointly evaluated datasets, reported in Figure 3 and Table 7. Because the comparison spans many benchmark tasks, we treat individual task wins as descriptive and base inferential statements on these paired aggregate differences; the bootstrap intervals quantify uncertainty in the mean task-level difference under task resampling, not variation across pretraining runs. Table 2 reports mean ROC-AUC by task group; per-dataset results for all models are given in Table 6.

Model	TDC-ADME	TDC-Tox	TDC-HTS	MoleculeNet	Overall
ECFP4	81.1	85.0	68.6	74.8	78.9
ChemBERTa-2 (MLM)	77.3	73.1	66.3	67.2	72.9
SELFormer	75.1	74.6	61.1	68.8	72.1
MoLFormer	81.5	85.2	66.8	77.6	79.8
ModernMolBERT-small	79.6	83.7	64.8	73.7	77.4
ModernMolBERT-base	80.7	83.8	64.4	73.7	77.9

Table 2: Mean ROC-AUC ($\times 100$) on the 25-dataset benchmark of Praski et al. [19], broken down by task group. Each entry averages per-dataset ROC-AUC using the best cross-validated downstream head (ridge / random forest / k NN); for multi-endpoint datasets, endpoint scores are averaged before aggregation. *Overall* is the unweighted mean across all 25 benchmark datasets. **Bold** marks the best value per column.

MODERNMOLBERT-base attains a mean ROC-AUC of 77.9 across all 25 benchmark datasets, and MODERNMOLBERT-small attains 77.4 across all 25 datasets (per-dataset results in Table 6). For comparison, ECFP4 reaches 78.9, MoLFormer 79.8, ChemBERTa-2 72.9, and SELFormer 72.1. Both MODERNMOLBERT variants therefore land between the strong fixed-fingerprint and large-scale SMILES baselines (ECFP4, MoLFormer) and the closest-scale pre-trained string encoders (ChemBERTa-2, SELFormer), which they exceed by roughly five ROC-AUC points. Figure 4 shows the corresponding per-dataset paired comparisons against each baseline.




Figure 3: **Paired bootstrap confidence intervals for ModernMolBERT-base vs. each baseline ($n = 25$ datasets).** Each row shows the mean Δ ROC-AUC ($\times 100$) between MODERNMOLBERT-base and one baseline (positive = MODERNMOLBERT-base higher), with 95% bootstrap confidence intervals ($B = 10,000$ resamples) and the number of datasets on which MODERNMOLBERT-base scores higher. The dashed line marks zero. The intervals for SELFormer and ChemBERTa-2 lie entirely above zero, that for MolFormer entirely below, and that for ECFP4 spans zero. Numerical values are given in Table 7 and full per-dataset results in Table 6. Resampling is over benchmark tasks, preserving the paired model difference within each task; the intervals therefore quantify uncertainty in the mean task-level difference, not variation across pretraining runs.

By task group, MODERNMOLBERT is strongest on the property-driven TDC tasks: it is competitive with ECFP4 on ADME (MODERNMOLBERT-base 80.7 vs. 81.1) and toxicity (83.8 vs. 85.0), the endpoints most closely tied to physicochemical structure (Figure 5). Its weakest group is the two-task TDC-HTS set (64.4), where all neural embedders trail ECFP4. On MoleculeNet, MolFormer shows its clearest advantage (77.6 vs. MODERNMOLBERT-base 73.7), consistent with its much larger and more chemically diverse pretraining corpus. This pattern is consistent with frozen learned embeddings transferring best to tasks governed by global physicochemical character rather than rare substructure recognition, an interpretation in line with broader analyses of when learned molecular representations help [29, 30].

Relative to ECFP4, MODERNMOLBERT approaches but does not surpass the fingerprint baseline. MODERNMOLBERT-base exceeds ECFP4 on 7 of the 25 datasets, with a mean Δ ROC-AUC of -1.0 and a 95% bootstrap confidence interval of $[-2.3, +0.4]$ (Table 7); the interval spans zero, consistent with neither a reliable advantage nor a reliable deficit at this scale. MODERNMOLBERT-small wins on 6 of 25. This is consistent with the continued strength of circular fingerprints in low-data property prediction.

The comparison most diagnostic of MODERNMOLBERT’s design is SELFormer, the closest prior SELFIES encoder (a RoBERTa-era backbone on the same representation). Here MODERNMOLBERT improves clearly and consistently: MODERNMOLBERT-base exceeds SELFormer on 23 of 25 datasets (mean Δ ROC-AUC $+5.8$; 95% CI $[+4.2, +7.5]$; Table 7), and MODERNMOLBERT-small on 24 of 25. Because both models consume SELFIES inputs, this gap is consistent with the combined benefit of MODERNBERT-era architectural choices, APE tokenisation, and differences in the training recipe; matched single-axis ablations would be required to assign the gain to individual components. MODERNMOLBERT likewise exceeds the SMILES-based ChemBERTa-2 by 5.0 points on the 25 datasets (mean Δ $+5.0$;




Figure 5: **Task-level ROC-AUC distributions by benchmark group.** Each point is one benchmark task evaluated under the frozen-embedding protocol with the best cross-validated downstream head, grouped by task family and model; the horizontal mark is the unweighted group mean (TDC-ADME, 12 tasks; TDC-Tox, 4; TDC-HTS, 2; MoleculeNet, 7) for ECFP4, ChemBERTa-2 (MLM), SELFormer, MolFormer, and the two released MODERNMOLBERT variants. For multi-endpoint datasets, endpoint ROC-AUCs are averaged before task-level aggregation. Colour redundantly encodes the model already labelled on the x -axis and follows a colourblind-safe palette used consistently across figures. The plot is descriptive; paired aggregate differences and their uncertainty are quantified in Figure 3 and Table 7. Masking-strategy ablation variants are omitted here and reported in Table 6.

25 datasets), all within the dataset-to-dataset scatter. Notably, the validation-MLM ranking does not carry through to downstream performance: heteroatom-biased span masking is weakest on the validation objective yet matches MODERNMOLBERT-base downstream, while span masking is strongest on validation MLM yet marginally trails standard masking. Under the frozen-embedding protocol we therefore find no reliable downstream effect of masking strategy.

Frozen-embedding quality is therefore largely insensitive to masking strategy and to the small-to-base size increase. This supports releasing the simplest configuration (standard masking) as the default and is consistent with the larger gains over SELFormer (Section 5.1) arising from the design choices shared by all released models. The corresponding per-task paired scatters isolating each pretraining axis are provided in Figure 9 (Section C).

6 Discussion

MODERNMOLBERT set out to test whether a compact encoder-only model can provide useful frozen molecular embeddings from chemically valid SELFIES strings and a SELFIES-aware APE tokeniser. The results support this premise: under a strictly frozen-embedding protocol, MODERNMOLBERT-base reaches a mean ROC-AUC of 77.9, approaching the strong ECFP4 fingerprint baseline and improving on the closest prior SELFIES and SMILES encoders by roughly five points. The subsections below set the scope of this claim, locate the gain among MODERNMOLBERT’s design choices, bound what the frozen-feature evaluation can establish, and outline what follows from it.

6.1 Scope and interpretation

This question is deliberately narrow. Its scope differs from large generative chemical language models, multimodal assay–molecule models, and graph or 3D molecular encoders. The




Figure 6: **Chemical property gradients across the ModernMolBERT embedding space.** PaCMAP projection [39] of frozen MODERNMOLBERT molecular embeddings, coloured by selected physicochemical properties from ChEMBL 36. A random subsample of 100,000 ChEMBL 36 molecules is shown. Each point represents one molecule; colour indicates the corresponding property value: ALogP, the quantitative estimate of drug-likeness (QED), polar surface area, rotatable bonds, aromatic rings, and heavy atom count. PaCMAP is used solely for two-dimensional visualisation of the high-dimensional embeddings; the absolute axis coordinates are not chemically meaningful, and only relative neighbourhoods and broad spatial trends should be interpreted. All panels show the same two-dimensional projection, recoloured by each descriptor.




Figure 7: **Primary sweep of MLM masking probability, masking strategy, and learning rate.** Validation performance across MLM masking probabilities for the small MODERNMOLBERT (~ 34 M parameters) and base MODERNMOLBERT (~ 114 M parameters) models. **(a,b)** evaluation masked accuracy (higher is better); **(c,d)** evaluation loss (lower is better); the x -axis is the MLM masking probability throughout. For the small model, line colour denotes masking strategy (span, hetero-span, standard); for the base model, all runs use standard masking and line colour denotes peak learning rate (2×10^{-4} , 4×10^{-4} , 8×10^{-4}). Shaded bands give the min-max range across runs and the asterisk marks the run that diverged during training. Full heatmaps across all conditions are provided in Figure 10.

resulting comparison is therefore most informative when interpreted as a test of frozen, string-based molecular representation learning rather than as a claim about all molecular foundation models. MODERNMOLBERT is used here purely as an encoder, not as a generative model; we therefore do not sample or reconstruct molecules from it. We note, however, that because both inputs and the MLM head operate over the SELFIES vocabulary, any SELFIES sequence the model predicts decodes to a chemically valid molecule by construction, so a generative use of the architecture would inherit SELFIES’ validity guarantee. We leave this direction to future work.

Every number in this paper should be read as a *frozen-embedding* result. No encoder is fine-tuned downstream; each produces a fixed mean-pooled embedding that is passed to lightweight classifiers. This protocol isolates representation transferability and mirrors fingerprint-based workflows, where fixed molecular features are precomputed once and reused across many endpoints. The fine-tuned behaviour of MODERNMOLBERT is left to future work (Section 6.4).

6.2 Advances over prior string models

Within this frozen-feature setting, the clearest advance is over the closest prior SELFIES encoder: holding the molecular representation fixed, MODERNMOLBERT improves on SELFormer by roughly five ROC-AUC points and is similarly ahead of the closest-scale SMILES MLM, ChemBERTa-2 (Section 5.1). Because the SELFormer comparison shares the SELFIES representation, this gain is consistent with benefits from the MODERNBERT backbone and APE tokeniser, though matched single-axis ablations would be required to isolate each component’s contribution. We do not overstate the result: MODERNMOLBERT is competitive with, but does not surpass, the strong ECFP4 fingerprint baseline and trails MoLFormer, which is pre-trained on roughly 400× more molecules [10]. The contribution is therefore a compact, reproducible SELFIES encoder that approaches fixed-fingerprint quality while substantially improving on prior string-based language models at comparable scale.

A practical advantage is accessibility. MODERNMOLBERT is released in the standard Hugging Face Transformers format and runs in the current Python ecosystem without model-specific inference kernels, legacy Python versions, or specialised runtime dependencies. Once molecules are provided as SELFIES strings, frozen embeddings can be extracted with a few lines of standard Transformers code, making the model straightforward to incorporate into screening, retrieval, and benchmark pipelines. The standalone APE tokeniser follows the same interface.

ModernMolBERT therefore occupies a practical trade-off point rather than a pure performance extreme. Larger or more specialised models remain preferable when maximum benchmark accuracy is the only objective, but many molecular-embedding workflows require representations that are adequate, reproducible, compact, and easy to deploy. In this setting, ModernMolBERT is a reasonable default for frozen string-based molecular embedding workflows: it approaches ECFP4, improves over comparable pretrained string encoders, and exposes frozen embeddings through the standard Hugging Face Transformers interface without specialised inference infrastructure.

6.3 Limitations

MODERNMOLBERT is pre-trained only on ChEMBL 36, a corpus dominated by drug-like, Lipinski-compliant compounds. Its training distribution is therefore narrower than

MoLFormer’s 1.1B PubChem [40] and ZINC [41] molecules, so generalisation to natural products, agrochemicals, peptides, macrocycles, or other under-represented chemical spaces may be weaker. This limitation is less severe for the drug-like benchmarks evaluated here, but broader pretraining (Section 6.4) would address it. Pretraining also uses canonical-SMILES-derived SELFIES strings only, with no randomised SMILES augmentation or multiple-enumeration, which have improved representation diversity in prior molecular language models [21, 42].

Two further constraints are properties of the string-encoder design rather than defects in this model. Like all sequential string models, MODERNMOLBERT has no access to explicit 3D conformational information, which graph- and geometry-based models such as Uni-Mol [28] and GEM [43] capture at the cost of greater architectural complexity; and the SELFIES-adapted APE tokeniser forms merge candidates from sequence-adjacent primitives, so merge boundaries need not align with graph-level bond adjacency. Neither affects the SELFIES validity guarantee or the frozen-embedding results reported here.

The evaluation is deliberately confined to frozen embeddings on binary classification tasks scored by ROC-AUC, following Praski et al. [19]. This isolates representation quality as an off-the-shelf feature vector and mirrors fingerprint workflows, but it does not measure the maximum performance reachable by supervised fine-tuning, nor does it establish behaviour on regression, retrieval, generative-conditioning, or structure-based tasks. A model with weaker frozen embeddings may still fine-tune well, and frozen evaluation may favour representations with broader but shallower chemical coverage.

Each released checkpoint was pre-trained from a single seed. The paired bootstrap intervals therefore quantify task-level variation under resampling of benchmark tasks, not pretraining-run variance, and the reported ranking should be read as a benchmark estimate for these released checkpoints rather than a full distribution over training runs; repeated-seed pretraining (Section 6.4) would supply that distribution.

The external baseline comparisons test the MODERNMOLBERT design package as a whole, not isolated causes. The improvement over SELFormer is consistent with the combined effect of the MODERNBERT architecture, APE tokenisation, training recipe, and implementation differences, but matched single-axis ablations (Section 6.4) would be required to attribute the gain to any individual component.

Finally, the ChEMBL 36 corpus and several benchmark datasets derive from overlapping medicinal-chemistry sources, and we do not deduplicate benchmark structures against the pretraining set. Some benchmark molecules or close analogues may therefore appear during pretraining, which could modestly inflate absolute scores for all pretrained encoders; the relative comparisons are less affected, since every pretrained baseline draws on similar public chemical space, and a structure-level overlap analysis (Section 6.4) would turn this caveat into a measured sensitivity check. More generally, systematic comparisons show that representation-learning gains are context-dependent, varying with data regime, split strategy, and activity cliffs [29, 30]. These considerations bear especially on frozen-feature evaluation, where the representation must transfer without task-specific adaptation.

6.4 Future Work

Several extensions follow naturally from the present design, all within the MLM-only, string-encoder scope of this work. The most direct is broader pretraining beyond ChEMBL, including PubChem, ZINC, Enamine, or other sources that better cover natural products,

fragments, reagents, and non-drug-like chemistry. Pretraining MODERNMOLBERT on the same corpora used by the baselines (for example the PubChem/ZINC mixture behind MoLFormer) would further allow the architecture and tokeniser to be compared at matched data, separating the contribution of pretraining scale from that of model design. A closely related question is how MODERNMOLBERT scales: the present study fixes a compact small/base pair and a 30,000-step budget, and systematically increasing model size, corpus size, and training compute would establish whether the remaining gap to large-scale SMILES models such as MoLFormer narrows with scale. A second direction is matched single-axis ablations of the tokeniser and backbone choices (substituting APE with standard BPE or the MODERNBERT backbone with a RoBERTa-era encoder) to isolate each component’s contribution to the gains over SELFormer. A third is to explore SELFIES encoders as molecular indexing models for chemical search, analogue retrieval, and dataset auditing, where deterministic valid string encodings and compact frozen embeddings are operationally attractive.

7 Conclusion

We presented MODERNMOLBERT, a compact encoder-only transformer model for SELFIES molecular language modelling. The released small and base variants combine the MODERNBERT encoder architecture with a SELFIES-adapted APE tokeniser and are pre-trained from scratch on a curated ChEMBL 36 corpus using the masked language modelling objective.

On all 25 benchmark datasets evaluated under the frozen embedding protocol, MODERNMOLBERT-base achieves a mean ROC-AUC of 77.9. This is competitive with the strong ECFP4 fingerprint baseline (78.9) and trails the much larger MoLFormer (79.8), while clearly exceeding the closest prior SELFIES and SMILES string encoders, SELFormer (72.1) and ChemBERTa-2 (72.9), by roughly five ROC-AUC points. Internal ablations further show that downstream embedding quality is largely insensitive to masking strategy and to the small-to-base size increase, so the gains over prior string models are consistent with the shared SELFIES, APE, and MODERNBERT design choices. These results suggest that a modern SELFIES encoder with an atom-preserving APE tokeniser can approach fingerprint-level frozen-feature performance while substantially improving on earlier string-based language models at comparable scale.

The central motivation is practical: molecular string encoders should be valid by construction, tokenised in chemically meaningful units, and built on modern efficient transformer components rather than inherited BERT-era backbones. By evaluating MODERNMOLBERT as a frozen embedder against ECFP4 and representative pre-trained molecular language models, we position a SELFIES-based ModernBERT encoder as a reproducible and computationally manageable foundation for molecular representation learning. Broader pretraining corpora and matched single-axis ablations are the natural next steps toward closing the remaining gap to large-scale models.

Data and Code Availability

Training code and evaluation scripts are available at <https://github.com/HauserGroup/ModernMolBERT>. Model weights and the APE tokeniser are available on Hugging Face at <https://huggingface.co/HauserGroup>. ChEMBL 36 is publicly available at <https://www.ebi.ac.uk/chembl/>.

[ebi.ac.uk/chembl/](https://ebl.ac.uk/chembl/). Source data for each quantitative figure are provided in `source_data/` in the code repository, with one machine-readable file per figure.

CRediT authorship contribution statement

J.S.M. Writing – original draft, Methodology, Formal analysis, Conceptualization. S.A. Methodology. C.L.P. Visualization, Writing – review & editing. X.L. Methodology. A.S.H. Methodology, Funding acquisition, Conceptualization.

Competing Interests

The authors declare no competing interests.

Use of AI Tools

During manuscript preparation, the authors used a large language model to assist with language editing, restructuring, and clarity improvements. All AI-assisted text was reviewed and edited by the authors, who take full responsibility for the final manuscript. No AI tool was used to generate, analyse, or interpret the experimental results.

Acknowledgements

We thank Marin Matic for valuable input on the analysis and for comments on the manuscript.

References

- [1] David Rogers and Mathew Hahn. Extended-connectivity fingerprints. *Journal of Chemical Information and Modeling*, 50(5):742–754, 2010. doi: 10.1021/ci100050t.
- [2] Sabrina Jaeger, Simone Fulle, and Samo Turk. Mol2vec: Unsupervised Machine Learning Approach with Chemical Intuition. *Journal of Chemical Information and Modeling*, 58(1):27–35, 2018. doi: 10.1021/acs.jcim.7b00616. URL <https://doi.org/10.1021/acs.jcim.7b00616>.
- [3] Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1263–1272. PMLR, 2017. URL <https://proceedings.mlr.press/v70/gilmer17a.html>.
- [4] Yu Rong, Yatao Bian, Tingyang Xu, Weiyang Xie, Ying Wei, Wenbing Huang, and Junzhou Huang. Self-supervised graph transformer on large-scale molecular data. In *Advances in Neural Information Processing Systems*, volume 33, pages 12559–12571, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/94aef38441efa3380a3bed3faf1f9d5d-Abstract.html>.
- [5] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT 2019*, pages 4171–4186, 2019. doi: 10.18653/v1/N19-1423.

- [6] David Weininger. SMILES, a chemical language and information system. 1. introduction to methodology and encoding rules. *Journal of Chemical Information and Computer Sciences*, 28(1):31–36, 1988. doi: 10.1021/ci00057a005.
- [7] Seyone Chithrananda, Gabriel Grand, and Bharath Ramsundar. ChemBERTa: Large-scale self-supervised pretraining for molecular property prediction. *arXiv preprint arXiv:2010.09885*, 2020. URL <https://arxiv.org/abs/2010.09885>.
- [8] Benedek Fabian, Thomas Edlich, H  l  na Gaspar, Marwin Segler, Joshua Meyers, Marco Fiscato, and Mohamed Ahmed. Molecular representation learning with language models and domain-relevant auxiliary tasks. *arXiv preprint arXiv:2011.13230*, 2020. URL <https://arxiv.org/abs/2011.13230>. MolBERT.
- [9] Walid Ahmad, Elana Simon, Seyone Chithrananda, Gabriel Grand, and Bharath Ramsundar. ChemBERTa-2: Towards chemical foundation models. *arXiv preprint arXiv:2209.01712*, 2022. doi: 10.48550/arXiv.2209.01712. URL <https://arxiv.org/abs/2209.01712>.
- [10] Jerret Ross, Brian Belgodere, Vijil Chenthamarakshan, Inkit Padhi, Youssef Mroueh, and Payel Das. Large-scale chemical language representations capture molecular structure and properties. *Nature Machine Intelligence*, 4(12):1256–1264, 2022. doi: 10.1038/s42256-022-00580-7. MolFormer.
- [11] Atakan Y  ksel, Erva Ulusoy, Atabey   nl  , and Tunca DoĖan. SELFormer: Molecular representation learning via SELFIES language models. *Machine Learning: Science and Technology*, 4(2):025035, 2023. doi: 10.1088/2632-2153/acdb30.
- [12] Mario Krenn, Florian H  se, AkshatKumar Nigam, Pascal Friederich, and Al  n Aspuru-Guzik. Self-referencing embedded strings (SELFIES): A 100% robust molecular string representation. *Machine Learning: Science and Technology*, 1(4):045024, 2020. doi: 10.1088/2632-2153/aba947.
- [13] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. *arXiv preprint arXiv:1508.07909*, 2016. URL <https://arxiv.org/abs/1508.07909>. Byte Pair Encoding (BPE).
- [14] Miguelangel Leon, Yuriy Perezhohin, Fernando Peres, Ale   Popovi  , and Mauro Castelli. Comparing SMILES and SELFIES tokenization for enhanced chemical language modeling. *Scientific Reports*, 14(1):25016, 2024. ISSN 2045-2322. doi: 10.1038/s41598-024-76440-8. URL <https://www.nature.com/articles/s41598-024-76440-8>.
- [15] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019. URL <https://arxiv.org/abs/1907.11692>.
- [16] Benjamin Warner, Antoine Chaffin, Benjamin Clavi  , Orion Weller, Oskar Hallstr  m, Said Taghadouini, Alexis Gallagher, Raja Biswas, Faisal Ladhak, Tom Aarsen, Nathan Cooper, Griffin Adams, Jeremy Howard, and Iacopo Poli. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference. *arXiv preprint arXiv:2412.13663*, 2024. URL <https://arxiv.org/abs/2412.13663>.

- [17] Mandar Joshi, Danqi Chen, Yinhan Liu, Daniel S. Weld, Luke Zettlemoyer, and Omer Levy. SpanBERT: Improving pre-training by representing and predicting spans. *Transactions of the Association for Computational Linguistics*, 8:64–77, 2020. doi: 10.1162/tac1_a_00300.
- [18] Tianhao Peng, Yuchen Li, Xuhong Li, Jiang Bian, Zeke Xie, Ning Sui, Shahid Mumtaz, Yanwu Xu, Linghe Kong, and Haoyi Xiong. Pre-trained molecular language models with random functional group masking. *npj Artificial Intelligence*, 1(1):28, 2025. ISSN 3005-1460. doi: 10.1038/s44387-025-00029-3. URL <https://www.nature.com/articles/s44387-025-00029-3>.
- [19] Mateusz Praski, Jakub Adamczyk, and Wojciech Czech. Benchmarking pretrained molecular embedding models for molecular representation learning. *arXiv preprint arXiv:2508.06199*, 2025. doi: 10.48550/arXiv.2508.06199. URL <https://arxiv.org/abs/2508.06199>.
- [20] Ross Irwin, Spyridon Dimitriadis, Jiazhen He, and Esben Jannik Bjerrum. Chemformer: A pre-trained transformer for computational chemistry. *Machine Learning: Science and Technology*, 3(1):015022, 2022. ISSN 2632-2153. doi: 10.1088/2632-2153/ac3ffb. URL <https://iopscience.iop.org/article/10.1088/2632-2153/ac3ffb>.
- [21] Robin Winter, Floriane Montanari, Frank No  , and Djork-Arn   Clevert. Learning continuous and data-driven molecular descriptors by translating equivalent chemical representations. *Chemical Science*, 10(6):1692–1701, 2019. doi: 10.1039/C8SC04175J.
- [22] Riya Singh, Aryan Amit Barsainyan, Rida Irfan, Connor Joseph Amorin, Stewart He, Tony Davis, Arun Thiagarajan, Shiva Sankaran, Seyone Chithrananda, Walid Ahmad, Derek Jones, Kevin McLoughlin, Hyojin Kim, Anoushka Bhutani, Shreyas Vinaya Sathyanarayana, Venkat Viswanathan, Jonathan E. Allen, and Bharath Ramsundar. ChemBERTa-3: An Open Source Training Framework for Chemical Foundation Models. *ChemRxiv*, 2025. doi: 10.26434/chemrxiv-2025-4glrl-v2. URL <https://chemrxiv.org/doi/full/10.26434/chemrxiv-2025-4glrl-v2>.
- [23] Fabian P. Kr  ger, Nicklas   sterbacka, Mikhail Kabeshov, Ola Engkvist, and Igor Tetko. MolEncoder: Towards optimal masked language modeling for molecules. *Digital Discovery*, 4(12):3552–3566, 2025. ISSN 2635-098X. doi: 10.1039/D5DD00369E. URL <https://xlink.rsc.org/?DOI=D5DD00369E>.
- [24] Philipp Seidl, Andreu Vall, Sepp Hochreiter, and G  nter Klambauer. Enhancing activity prediction models in drug discovery with the ability to understand human language. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. URL <https://proceedings.mlr.press/v202/seidl23a.html>. CLAMP.
- [25] Benjamin Kaufman, Edward C. Williams, Carl Underkoffler, Ryan Pederson, Narbe Mardirossian, Ian Watson, and John Parkhill. COATI: Multimodal contrastive pretraining for representing and traversing chemical space. *Journal of Chemical Information and Modeling*, 64(4):1145–1157, 2024. doi: 10.1021/acs.jcim.3c01753.
- [26] Weihua Hu, Bowen Liu, Joseph Gomes, Marinka Zitnik, Percy Liang, Vijay Pande, and Jure Leskovec. Strategies for pre-training graph neural networks. In *International*

- Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1905.12265>.
- [27] Shengchao Liu, Hanchen Wang, Weiyang Liu, Jon Lasenby, Hongyu Guo, and Jian Tang. Pre-training molecular graph representation with 3D geometry. In *International Conference on Learning Representations*, 2022. URL <https://arxiv.org/abs/2110.07728>.
- [28] Gengmo Zhou, Zhifeng Gao, Qiankun Ding, Hang Zheng, Hongteng Xu, Zhewei Wei, Linfeng Zhang, and Guolin Ke. Uni-Mol: A universal 3D molecular representation learning framework. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=6K2RM6wVqKu>.
- [29] Kevin Yang, Kyle Swanson, Wengong Jin, Connor Coley, Philipp Eiden, Hua Gao, Angel Guzman-Perez, Timothy Hopper, Brian Kelley, Miriam Mathea, Andrew Palmer, Volker Settels, Tommi Jaakkola, Klavs Jensen, and Regina Barzilay. Analyzing learned molecular representations for property prediction. *Journal of Chemical Information and Modeling*, 59(8):3370–3388, 2019. ISSN 1549-9596. doi: 10.1021/acs.jcim.9b00237.
- [30] Jianyuan Deng, Zhibo Yang, Hehe Wang, Iwao Ojima, Dimitris Samaras, and Fusheng Wang. A systematic study of key elements underlying molecular property prediction. *Nature Communications*, 14(1):6395, 2023. ISSN 2041-1723. doi: 10.1038/s41467-023-41948-6. URL <https://www.nature.com/articles/s41467-023-41948-6>.
- [31] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. doi: 10.1016/j.neucom.2023.127063.
- [32] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. *arXiv preprint arXiv:2205.14135*, 2022. URL <https://arxiv.org/abs/2205.14135>.
- [33] Mike Mayuare. APE tokenizer: Atom pair encoding for chemical strings. GitHub repository, 2024. URL <https://github.com/mikemayuare/apetokenizer>.
- [34] Benjamin Warner, Benjamin Clavié, Antoine Chaffin, Oskar Hallström, et al. Finally, a replacement for BERT: Introducing ModernBERT. Hugging Face Blog, 2024. URL <https://huggingface.co/blog/modernbert>. Published 19 December 2024.
- [35] Barbara Zdrazil, Eloy Felix, Fiona Hunter, Emma J Manners, James Blackshaw, Sybilla Corbett, Marleen de Veij, Harris Ioannidis, David Mendez Lopez, Juan F Mosquera, Maria Paula Magarinos, Nicolas Bosc, Ricardo Arcila, Tevfik Kizilören, Anna Gaulton, A Patrícia Bento, Melissa F Adasme, Peter Monecke, Gregory A Landrum, and Andrew R Leach. The ChEMBL database in 2023: A drug discovery platform spanning multiple bioactivity data types and time periods. *Nucleic Acids Research*, 52(D1):D1180–D1192, 2024. ISSN 0305-1048. doi: 10.1093/nar/gkad1004. URL <https://academic.oup.com/nar/article/52/D1/D1180/7337608>.
- [36] Greg Landrum. RDKit: Open-source cheminformatics. Software, 2023. URL <https://www.rdkit.org>. Version 2023.09.

- [37] Kexin Huang, Tianfan Fu, Wenhao Gao, Yue Zhao, Yusuf Roohani, Jure Leskovec, Connor W Coley, Cao Xiao, Jimeng Sun, and Marinka Zitnik. Therapeutics Data Commons: Machine Learning Datasets and Tasks for Drug Discovery and Development. *arXiv preprint arXiv:2102.09548*, 2021. URL <https://tdcommons.ai>.
- [38] Zhenqin Wu, Bharath Ramsundar, Evan N. Feinberg, Joseph Gomes, Caleb Geniesse, Aneesh S. Pappu, Karl Leswing, and Vijay Pande. MoleculeNet: A benchmark for molecular machine learning. *Chemical Science*, 9(2):513–530, 2018. doi: 10.1039/C7SC02664A.
- [39] Yingfan Wang, Haiyang Huang, Cynthia Rudin, and Yaron Shaposhnik. Understanding How Dimension Reduction Tools Work: An Empirical Approach to Deciphering t-SNE, UMAP, TriMap, and PaCMAP for Data Visualization. *Journal of Machine Learning Research*, 22(201):1–73, 2021. URL <https://jmlr.org/papers/v22/20-1061.html>.
- [40] Sunghwan Kim, Jie Chen, Tiejun Cheng, Asta Gindulyte, Jia He, Siqian He, Qingliang Li, Benjamin A. Shoemaker, Paul A. Thiessen, Bo Yu, Leonid Zaslavsky, Jian Zhang, and Evan E. Bolton. PubChem 2023 update. *Nucleic Acids Research*, 51(D1):D1373–D1380, 2023. doi: 10.1093/nar/gkac956.
- [41] John J. Irwin, Khanh G. Tang, Jerome Young, Chinzorig Dandarchuluun, Benjamin R. Wong, Munkhzul Khurelbaatar, Yurii S. Moroz, John Mayfield, and Roger A. Sayle. ZINC20 — a free ultralarge chemical database for ligand discovery. *Journal of Chemical Information and Modeling*, 60(12):6065–6073, 2020. doi: 10.1021/acs.jcim.0c00675.
- [42] Chae Eun Lee, Jin Sob Kim, Jin Hong Min, and Sung Won Han. SimSon: Simple contrastive learning of SMILES for molecular property prediction. *Bioinformatics*, 41(5):btaf275, 2025. doi: 10.1093/bioinformatics/btaf275. URL <https://academic.oup.com/bioinformatics/article/41/5/btaf275/8127203>.
- [43] Xiaomin Fang, Lihang Liu, Jieqiong Lei, Donglong He, Shanzhuo Zhang, Jingbo Zhou, Fanzhong Wang, Hua Wu, and Haifeng Wang. Geometry-enhanced molecular representation learning for property prediction. *Nature Machine Intelligence*, 4:127–134, 2022. doi: 10.1038/s42256-021-00438-4.

Appendix A Training Hyper-parameters

Hyper-parameter	ModernMolBERT-small	ModernMolBERT-base
Parameters	34.1M	114.3M
Hidden size	512	768
Transformer layers	8	12
Attention heads	8	12
FFN intermediate size	2048	3072
Global attn. every n layers	3	3
Local attention window	128	128

Table 3: Hyper-parameters for the two MODERNMOLBERT variants. Both use the MODERN-BERT backbone with alternating local/global attention (global every 3 layers; local window of 128 tokens).

Hyperparameter	Small	Base
<i>Fixed</i>		
Max training steps	30,000	30,000
Max sequence length	128	128
Effective batch size	256	256
Per-device batch size	256	64
Gradient accumulation	1	4
Weight decay	0.01	0.01
Gradient clip norm	1.0	1.0
Precision	bf16	bf16
Seed	42	42
<i>Swept (3×3 grid, 9 runs each)</i>		
MLM masking rate	{0.15, 0.20, 0.25}	
Peak learning rate	$\{10^{-4}, 2 \times 10^{-4}, 4 \times 10^{-4}\}$	$\{2 \times 10^{-4}, 4 \times 10^{-4}, 8 \times 10^{-4}\}$
LR warmup steps	1,500	2,000

Table 4: Pretraining hyperparameters for both MODERNMOLBERT variants. Swept values are shown in braces. The released checkpoints use the stable standard-masking configuration (4×10^{-4} for small, 2×10^{-4} for base); the full sweep is analysed in Section 5.2.




Figure 8: **Pretraining convergence of the released small ModernMolBERT.** Validation curves for the small model (standard masking, 4×10^{-4}) over 30,000 training steps. **(a)** Validation masked-token accuracy. **(b)** Training loss (dense) and validation MLM loss (markers). Both curves plateau smoothly, confirming stable convergence within the fixed 30,000-step budget. Per-step logs for the base model were not retained; its final validation metrics are reported in the sweep (Figure 7).

Appendix B Tokeniser Details

Tokeniser property	Value
<i>Vocabulary</i>	
Training molecules	2,000,000
Base vocabulary (APE merges)	589
Added SELFIES primitives	42
Final vocabulary	631
<i>Sequence statistics (10 000 held-out ChEMBL 36 molecules)</i>	
Median length (tokens)	24
Mean length (tokens)	25.7
95th percentile	43
99th percentile	56
<i>Coverage</i>	
Truncation at 128 tokens	<0.1%
Unknown-token rate	0.0%

Table 5: Summary statistics for the SELFIES-adapted APE tokeniser. The base vocabulary is learned by APE merge training on 2M randomly sampled ChEMBL 36 SELFIES strings; 42 additional SELFIES primitives observed in downstream benchmark molecules are appended after merge learning but before model pretraining, giving a final vocabulary of 631 tokens on which all model weights are initialised and trained. Sequence statistics and coverage are evaluated on 10 000 held-out ChEMBL 36 molecules.

The final tokeniser is trained on 2M randomly sampled ChEMBL 36 SELFIES strings with `max_merge_pieces=2` and `min_freq_for_merge=3000`, yielding a base vocabulary of 589 tokens. An additional 42 SELFIES primitives observed in the benchmark datasets (minimum frequency ≥ 10) are appended after APE merge learning but before model pretraining, giving a final vocabulary of 631 tokens on which all model weights were initialised and trained from the outset. Special tokens are assigned fixed ids: [BOS] = 0, [PAD] = 1, [EOS] = 2, [UNK] = 3, [MASK] = 4. Evaluated on 10 000 held-out ChEMBL 36 molecules, the tokeniser achieves

an unknown-token rate of 0.0%, a median sequence length of 24 tokens (mean 25.7; 95th percentile 43; 99th percentile 56), negligible truncation at the model context length of 128 tokens (<0.1%), and zero truncations at 256 tokens.

Appendix C Additional Benchmark Results

This appendix reports per-task benchmark results, together with auxiliary analyses used to assess robustness of the frozen-feature evaluation. Table 6 gives the full per-task test ROC-AUC for every model, including the span- and heteroatom-span masking variants of MODERNMOLBERT-small. All four MODERNMOLBERT checkpoints and the four baselines are evaluated on the full set of 25 datasets.

Task	ECFP4	ChBa-2	SELF.	MoLF.	MMB-s	MMB-b	MMB-sp	MMB-h
<i>TDC – ADME</i>								
Bioavailability	65.3	65.2	64.1	72.2	63.1	65.1	67.6	71.1
HIA	93.9	88.7	88.4	98.1	97.8	98.0	98.1	97.9
Pgp	91.6	84.3	83.3	92.9	91.5	93.0	90.6	92.0
PAMPA	77.4	67.9	70.8	71.6	73.9	75.5	72.2	75.5
CYP1A2	91.9	87.6	86.5	91.4	90.4	91.3	89.9	90.9
CYP2C19	86.8	81.8	76.5	85.5	84.2	84.9	83.1	84.4
CYP2C9	86.7	82.1	80.5	86.6	84.9	84.9	84.6	85.3
CYP2D6	86.3	78.2	79.1	85.5	83.3	85.2	83.6	85.1
CYP3A4	86.3	80.5	78.7	85.0	83.1	83.6	81.9	83.3
CYP2C9 (substrate)	62.1	65.6	60.4	62.5	62.0	62.6	60.7	59.4
CYP2D6 (substrate)	79.0	78.7	71.4	83.0	79.2	82.0	78.8	78.8
CYP3A4 (substrate)	65.4	67.2	61.2	63.1	62.1	62.4	65.0	66.9
<i>TDC – Toxicity</i>								
AMES	84.7	74.7	70.5	81.7	78.3	78.2	78.2	78.2
DILI	88.2	73.9	74.3	92.0	92.3	89.3	93.2	90.2
hERG	80.2	64.3	79.1	83.1	81.6	84.9	80.5	80.2
hERG (Karim)	87.1	79.5	74.5	84.1	82.4	82.8	81.4	82.2
<i>TDC – HTS</i>								
SARS-CoV-2 3CLPro	76.1	67.6	66.6	74.9	70.7	73.2	73.0	71.9
SARS-CoV-2 (Vitro)	61.0	64.9	55.7	58.7	58.9	55.6	54.6	56.5
<i>MoleculeNet</i>								
BACE	86.2	80.8	78.4	83.7	78.4	81.5	78.6	80.1
BBBP	70.2	70.7	65.4	71.5	69.3	68.3	69.3	68.3
ClinTox	73.1	55.9	79.0	88.4	79.2	82.1	81.6	83.2
HIV	78.4	75.3	75.9	78.0	76.2	73.3	74.3	75.5
MUV	72.4	56.6	57.7	79.8	74.0	72.1	69.3	72.6
SIDER	67.4	61.6	56.6	66.6	64.7	64.3	63.5	66.0
Tox21	75.9	69.4	69.0	75.0	73.9	74.2	73.6	72.8

Table 6: Per-task test ROC-AUC ($\times 100$) for all models on the 25-task benchmark. *MMB-s* = MODERNMOLBERT-small (standard), *MMB-b* = MODERNMOLBERT-base, *MMB-sp* = small span masking, *MMB-h* = small hetero-span masking. All models are evaluated on the full set of 25 tasks.

Model A	Model B	n	Wins	Ties	Losses	Mean Δ	95 % CI
MMB-base	SELFormer	25	23	0	2	+5.79	[+4.2, +7.5]
MMB-base	ChemBERTa-2	25	19	0	6	+5.00	[+2.1, +8.1]
MMB-base	ECFP4	25	7	0	18	-1.02	[-2.3, +0.4]
MMB-base	MoLFormer	25	4	0	21	-1.87	[-2.9, -0.9]

Table 7: Paired bootstrap confidence intervals (95 %, $B = 10,000$ resamples) on mean Δ ROC-AUC ($\times 100$) between MODERNMOLBERT-base and each baseline, computed over the tasks where both models have a result. *Wins/Ties/Losses* count tasks where MODERNMOLBERT-base is above, equal to, or below the baseline. A positive mean Δ and CI entirely above zero indicates a consistent advantage for MODERNMOLBERT-base.




Figure 9: **Internal comparison of the four ModernMolBERT checkpoints on downstream ROC-AUC.** Paired scatter of held-out test ROC-AUC across the benchmark tasks (one point per task, best cross-validated downstream head), isolating each pretraining design axis against the released small model (standard masking, y -axis throughout except panel a). **(a)** MODERNMOLBERT-small versus MODERNMOLBERT-base, isolating model size. **(b)** MODERNMOLBERT-small versus the span-masking variant. **(c)** MODERNMOLBERT-small versus the heteroatom-biased span-masking variant. The dashed line is the identity $y = x$, points are coloured by task group, and each inset reports the number of tasks on which the y -axis model wins. All three panels are evaluated on the full 25 datasets (see Table 6). These ablations vary masking strategy and model size within the MODERNMOLBERT setup; they do not isolate the effects of SELFIES, APE tokenisation, or the MODERNBERT backbone relative to matched alternatives.

Appendix D Molecular Embedding Models

Table 8 lists the non-graph molecular embedding models evaluated in the benchmark of Praski et al. [19]. Graph-based models (R-MAT, MAT, GROVER, GIN, MolR, GraphMVP, Uni-Mol, Uni-Mol2, GEM, GraphFP) are excluded here as they fall outside the scope of MODERNMOLBERT. Table 6 of Praski et al. [19] provides the full benchmark results across all 25 methods, including graph-based architectures.

Model	Input	Params (M)	Dim.	Year	Corpus	Size (M)
Atom Pair	Fingerprint	—	2,048	1985	—	—
TT	Fingerprint	—	2,048	1987	—	—
ECFP	Fingerprint	—	2,048	2010	—	—
Mol2Vec	Substructure	n.r.	300	2018	ChEMBL	n.r.
CDDD	SMILES	n.r.	512	2019	n.r.	n.r.
MolBERT	SMILES	85	768	2020	ChEMBL	2
ChemBERTa	SMILES	5–46 ^a	384	2020	PubChem	77
ChemGPT	SMILES/SELFIES	5	128	2022	PubChem	10
Chemformer	SMILES	45	512	2022	ZINC-15	100
MolFormer	SMILES	46.8	768	2022	PubChem + ZINC	>1,100
CLAMP	Fingerprint/RDKit	n.r.	768	2023	PubChem assays	143 meas.
COATI	SMILES/3D	n.r.	256	2023	n.r.	n.r.
SELFormer	SELFIES	87	768	2023	ChEMBL	2
SimSon	SMILES	n.r.	512	2025	n.r.	n.r.
ChemFM	SMILES	3,000	3,072	2025	UniChem	178
ModernMolBERT	SELFIES	34/114	512/768	2026	ChEMBL 36	2.4

Table 8: Non-graph molecular embedding models considered in this work, following the benchmark selection of Praski et al. [19]. Graph-based architectures are excluded; see Table 6 of that work for the full model comparison. *Params*: trainable parameters in millions; *Size*: pretraining corpus size in millions unless otherwise stated; n.r. = not reported clearly in the paper, model card, or repository; “—” = not applicable. ^aFor ChemBERTa, 5M/10M/77M denote pretraining corpus sizes, not parameter counts; reported architecture searches span approximately 5–46M parameters.

Appendix E Pretraining Hyperparameter Heatmaps

This appendix reports the full hyperparameter sweep underlying the masking and learning-rate selections discussed in Section 5.2, as heatmaps over all masking probabilities, strategies, and learning rates evaluated.




Figure 10: Masked language modelling sweep across masking strategies, masking probabilities, and learning rates. Heatmaps show validation performance for the small MODERNMOLBERT model (~ 34 M parameters) and the base MODERNMOLBERT model (~ 114 M parameters). Panels e and f show masked-token accuracy (higher is better). Panels g and h show evaluation loss (lower is better). For the small model, columns compare masking strategies: span masking, hetero-span masking, and standard token masking. For the base model, columns compare peak learning rates under standard masking: 2×10^{-4} , 4×10^{-4} , and 8×10^{-4} . Rows give the MLM masking probability: 0.15, 0.20, and 0.25. For the small model, the second line in each cell reports the learning rate used; asterisks denote runs that diverged during training. On validation MLM metrics, the small model performs best with span masking at a learning rate of 2×10^{-4} , with accuracy increasing slightly as masking probability increases. Hetero-span masking is consistently worse for both accuracy and loss. For the base model, 2×10^{-4} and 4×10^{-4} perform similarly; 8×10^{-4} produces unstable training with superficially moderate final accuracy but elevated loss. Line-plot views of panels a–d are shown in the main text (Figure 7).